

Random Forests and Boosting

Travers Parsons-Grayson

April 7, 2019

Random Forests vs. Bagging (Motivation)

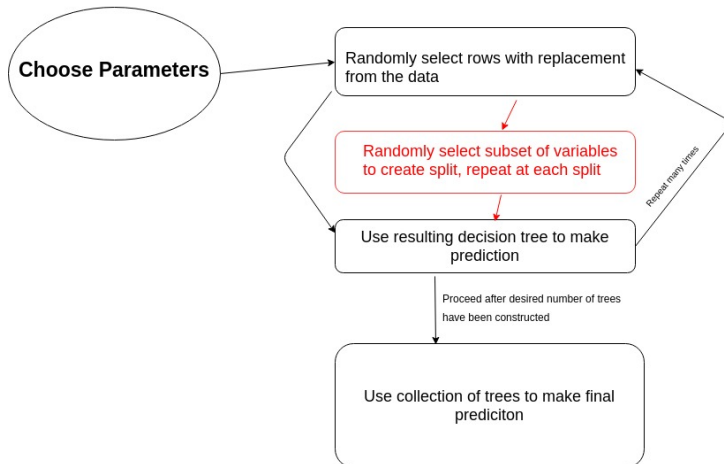
The Difference

The process of creating a random forest is very similar to the process of bagging with one small caveat. In a random forest every time a split is considered a random sample of m predictors from the total p predictors are chosen as candidates for the split. Bagging is a special case of Random Forests when $m = p$.

Why?

Decorrelation: The weakness of bagging is that many trees end up looking the same because they will almost always use the strongest predictors in the same order. Random Forests *decorrelate* trees and thus reduce the variance of the prediction.

Bagging vs Random Forests



Algorithm

Parameters

- ① Number of trees, k
- ② Number of variables to select randomly at each split, m
- ③ (optional) Size of training set, the rows that we sample from without replacement
- ④ (optional) Maximum size of the trees grown, by number of nodes j

Construction

- ① Randomly select rows with replacement from data (typically use 2/3's of rows)

Algorithm

Parameters

- ① Number of trees, k
- ② Number of variables to select randomly at each split, m
- ③ (optional) Size of training set, the rows that we sample from without replacement
- ④ (optional) Maximum size of the trees grown, by number of nodes j

Construction

- ① Randomly select rows with replacement from data (typically use 2/3's of rows)
- ② **Randomly select m variables to create split** (typically $m \equiv \sqrt{p}$)

Algorithm

Parameters

- ① Number of trees, k
- ② Number of variables to select randomly at each split, m
- ③ (optional) Size of training set, the rows that we sample from without replacement
- ④ (optional) Maximum size of the trees grown, by number of nodes j

Construction

- ① Randomly select rows with replacement from data (typically use 2/3's of rows)
- ② **Randomly select m variables to create split** (typically $m \equiv \sqrt{p}$)
- ③ Repeat step 2 at each split until decision tree is built

Algorithm

Parameters

- ① Number of trees, k
- ② Number of variables to select randomly at each split, m
- ③ (optional) Size of training set, the rows that we sample from without replacement
- ④ (optional) Maximum size of the trees grown, by number of nodes j

Construction

- ① Randomly select rows with replacement from data (typically use 2/3's of rows)
- ② **Randomly select m variables to create split** (typically $m \equiv \sqrt{p}$)
- ③ Repeat step 2 at each split until decision tree is built
- ④ Use resulting decision tree to make prediction

Algorithm

Parameters

- ① Number of trees, k
- ② Number of variables to select randomly at each split, m
- ③ (optional) Size of training set, the rows that we sample from without replacement
- ④ (optional) Maximum size of the trees grown, by number of nodes j

Construction

- ① Randomly select rows with replacement from data (typically use 2/3's of rows)
- ② **Randomly select m variables to create split** (typically $m \equiv \sqrt{p}$)
- ③ Repeat step 2 at each split until decision tree is built
- ④ Use resulting decision tree to make prediction
- ⑤ Repeat steps 1-3 k times

Example in R (Ames Housing)

We will use a random forest to predict the sale price of a house in the Ames Housing dataset.

```
# We will use the randomForest library  
library(randomForest)  
  
# Use a selection of variables  
rf <- randomForest(Sales_Price ~ Year_Sold + Alley  
                    + Year_Built, data = ames)  
  
# Use all variables  
rf <- randomForest(Sales_Price ~ ., data = ames)
```

- We will use the randomForest library in R to construct our randomForests
- Simply provide the constructor with a formula and the dataset to be used

Variable Selection

```
```{r, warning=FALSE, message=FALSE,echo = FALSE,tidy = TRUE}
library(readr)

=====
Commented because running script is very
computationally expensive
=====

oob.err=double(18)
test.err=double(18)

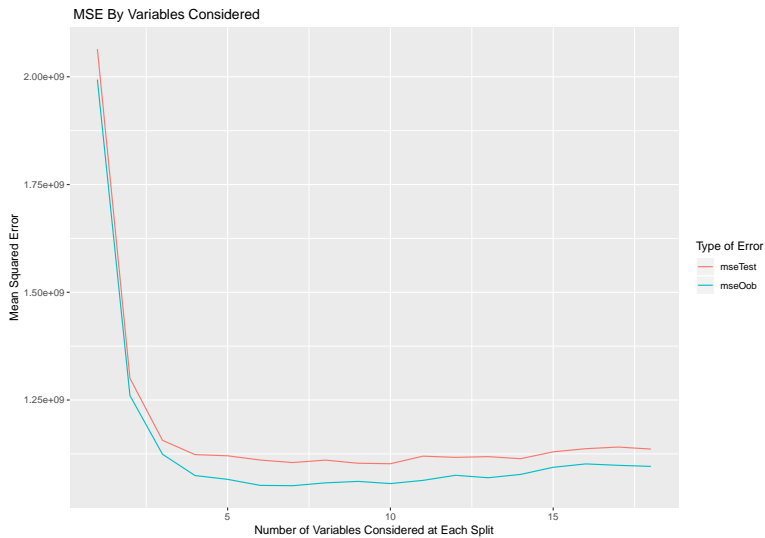
Try building a random forest for every value of mtry
for(mtry in 1:18)
{
 rf= randomForest(Sale_Price ~ ., data = train,mtry = mtry)
 oob.err[mtry] = rf$mse[400] #Error of all Trees fitted

 pred<-predict(rf,test) #Predictions on Test Set for each Tree
 test.err[mtry]= with(test, mean((Sale_Price - pred)^2)) #Mean Squared Test Error

 cat(mtry," ") #printing the output to the console
}
errFrame = data.frame(numVars = 1:18,mseTest = test.err,mseOob = oob.err)
write.csv(errFrame,"error_by_split.csv")

errFrame <- read_csv("error_by_split.csv")

errFrame <- reshape2::melt(errFrame,id = 'numVars')
```



# Creating the Random Forest

```
{r, load myData, warning=FALSE, message=FALSE}
library(AmesHousing)
library(dplyr)
library(ggplot2)
library(randomForest)

load("myData.RData")

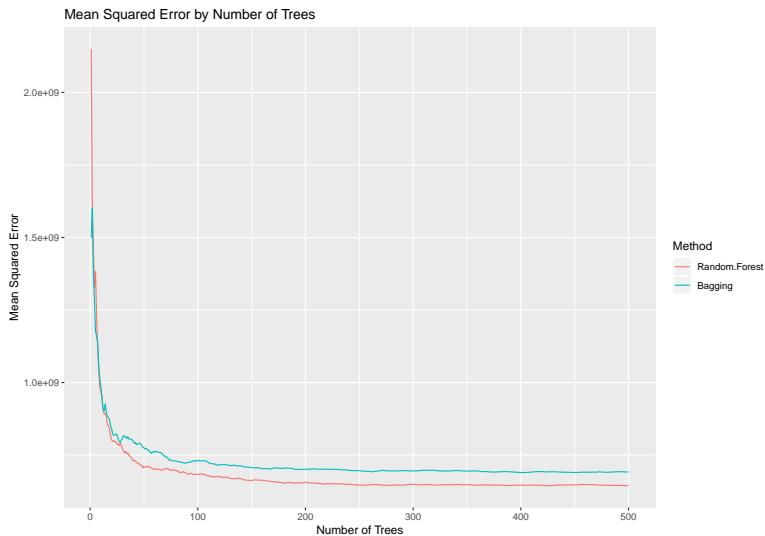
set.seed(15)

Reduce the number of columns from 81 to 18
ames <- make_ames()
ames <- ames %>% select(c("Sale_Type", "Pool_QC", "Year_Sold", "Alley", "Overall_Cond", "Year_Built", "Lot_Area", "Street",
 "Central_Air", "Lot_Frontage", "MS_Zoning", "House_Style", "Bldg_Type", "Neighborhood",
 "Overall_Qual", "Utilities", "Exter_Qual", "Sale_Price"))

Split the data into training and test set
sample <- sample.int(n = nrow(ames), size = floor(.8*nrow(ames)), replace = F)
train <- ames[sample,]
test <- ames[-sample,]

Build a random forest and Bagging
amesRF <- randomForest(Sale_Price ~ ., data = train, mtry = 10)
amesBag <- randomForest(Sale_Price ~ ., data = train, mtry = 18)
```





# Boosting

Boosting differs from Random Forests in that instead of growing many *decorrelated* trees, we build trees sequentially using information from previous trees.

## Main Idea

Fit the first tree using the actual outcomes. Then from there on out use residuals from previous tree to fit next tree.

# Algorithm

## Parameters

- ①  $B$ , the number of trees to build. Unlike random forests there is a risk of overfitting the model if  $B$  is too high. Use cross-validation to find optimal value of  $B$ .
- ②  $\lambda$ , the shrinkage factor (typically 0.01 or 0.0001). The shrinkage factor determines how much the residuals from the current decision tree change the overall model.
- ③  $d$ , the number of splits in each tree. When  $d = 1$ , each tree is called a *stump*.



## Construction

Let  $\hat{f}$  be our function with which we will make regression estimates.

- ① Let  $\hat{f}(x) = 0$  and  $r_i = y_i$  for all  $i$ .
- ② For  $b \in \{1, 2, \dots, B\}$ , repeat:
  - ① Fit a tree  $t_b$  with  $d$  splits to the training data ( $X$  and  $r=y$ ).
  - ② Update  $\hat{f}(x)$  by adding weighted version of new tree:

$$\hat{f}(x) = \hat{f}(x) + \lambda t_b(x).$$

- ③ Update the residuals  $r$ :

$$r_i = r_i - \lambda t_b(x_i).$$

## Construction

Let  $\hat{f}$  be our function with which we will make regression estimates.

- ① Let  $\hat{f}(x) = 0$  and  $r_i = y_i$  for all  $i$ .
- ② For  $b \in \{1, 2, \dots, B\}$ , repeat:
  - ① Fit a tree  $t_b$  with  $d$  splits to the training data ( $X$  and  $r=y$ ).
  - ② Update  $\hat{f}(x)$  by adding weighted version of new tree:

$$\hat{f}(x) = \hat{f}(x) + \lambda t_b(x).$$

- ③ Update the residuals  $r$ :

$$r_i = r_i - \lambda t_b(x_i).$$

- ③ Return the final model:

$$\hat{f}(x) = \sum_{b=1}^B \lambda t_b(x)$$

## Example in R

For boosting, we will use the *gbm* package. Once again we will use the Ames Housing Data.

```
library(gbm)
set.seed(1)
Create our boosting model
amesBoost <- gbm(Sale_Price ~ ., data = ames, train.fraction = .7)
```

- Formula and data parameters work similarly to `randomForest()`
- *train.fraction* determines the fraction of data to be used for training and the fraction to be used for validation
- *distribution*, typically *gaussian* if outcome variable is continuous and *bernoulli* for binary outcome
- *n.trees*, number of trees to train model on
- *interaction.depth*, maximum number of splits allowed in each tree

